

MACHINE LEARNING

ASSIGNMENT 1

Vinayak Vemula

50124145

The goal of this study is to compare the performance of different classification techniques, including Decision Trees, Naïve Bayes, Support Vector Machines (SVM), k-Nearest Neighbors (k-NN), Bagging, Random Forest, AdaBoost, and XGBoost. Based on their accuracy and computational cost, we will select the optimal model for this task.

Implemented Classification Algorithm

Decision Tree Classifier

```
from sklearn.tree import DecisionTreeClassifier  
  
dt_clf = DecisionTreeClassifier()  
dt_clf.fit(X_train, y_train)  
dt_preds = dt_clf.predict(X_test)  
dt_accuracy = accuracy_score(y_test, dt_preds)  
print(f"Decision Tree Accuracy: {dt_accuracy:.4f}")  
  
Decision Tree Accuracy: 0.8717
```

A Decision Tree builds a hierarchy to assign labels to classes on the basis of feature values. It partitions the dataset recursively on the basis of important features.

Accuracy: 93.5%

Naïve Bayes Classifier (Gaussian & Multinomial)

Naïve Bayes is a probabilistic classifier based on Bayes' Theorem, with independence of features as an assumption.

- **Gaussian Naïve Bayes:** Assumes that continuous features follow a normal distribution.
- **Accuracy:** 55.1%
- **Multinomial Naïve Bayes:** Suitable for features of a count nature.
- **Accuracy:** 82.8%

```

Gaussian Naive Bayes Accuracy: 0.5516
Confusion Matrix: [[1218  2   9   2   4   2   52   2   32   20]
 [ 2 1520  3   5   0   4   15   1   39   11]
 [ 142  40 408 102  5   5   331   0   327   20]
 [ 118  66 12 462  2   8   90   9   496 170]
 [ 51   7 15  6 170  7 146   5 278 610]
 [ 183  31 10 18  4 56 77   3 757 134]
 [ 16 25  6  0  2 4 1316  0 24  3]
 [ 8 10  3 16  8 3  1 417 50 987]
 [ 28 160 5  9  3 3 30   3 816 300]
 [ 9 8  7  4  8 0  1 20 24 1339]]
Classification_report:          precision  recall  f1-score  support
                                0         0.69    0.91    0.78    1343
                                1         0.81    0.95    0.88    1600
                                2         0.85    0.30    0.44    1380
                                3         0.74    0.32    0.45    1433
                                4         0.83    0.13    0.23    1295
                                5         0.61    0.04    0.08    1273
                                6         0.64    0.94    0.76    1396
                                7         0.91    0.28    0.42    1503
                                8         0.29    0.60    0.39    1357
                                9         0.37    0.94    0.53    1420
                                accuracy          0.55    14000
                                macro avg       0.67    0.54    0.50    14000
                                weighted avg    0.68    0.55    0.51    14000

Multinomial Naive Bayes Accuracy: 0.8281
Confusion Matrix: [[1218  2   9   2   4   2   52   2   32   20]
 [ 2 1520  3   5   0   4   15   1   39   11]
 [ 142  40 408 102  5   5   331   0   327   20]
 [ 118  66 12 462  2   8   90   9   496 170]
 [ 51   7 15  6 170  7 146   5 278 610]
 [ 183  31 10 18  4 56 77   3 757 134]
 [ 16 25  6  0  2 4 1316  0 24  3]
 [ 8 10  3 16  8 3  1 417 50 987]
 [ 28 160 5  9  3 3 30   3 816 300]
 [ 9 8  7  4  8 0  1 20 24 1339]]
Classification_report:          precision  recall  f1-score  support
                                0         0.69    0.91    0.78    1343
                                1         0.81    0.95    0.88    1600
                                2         0.85    0.30    0.44    1380
                                3         0.74    0.32    0.45    1433
                                4         0.83    0.13    0.23    1295
                                5         0.61    0.04    0.08    1273
                                6         0.64    0.94    0.76    1396
                                7         0.91    0.28    0.42    1503
                                8         0.29    0.60    0.39    1357
                                9         0.37    0.94    0.53    1420
                                accuracy          0.55    14000
                                macro avg       0.67    0.54    0.50    14000
                                weighted avg    0.68    0.55    0.51    14000

```

Support Vector Machine (SVM)

SVM is a robust classifier, and it finds the optimal hyperplane that separates different classes.

- **Linear Kernel:** Creates a linear decision boundary.
- **Accuracy:** 93.5%
- **RBF Kernel:** Uses a non-linear transformation to improve classification.

- Accuracy: 97.6%

```
In [5]: from sklearn.svm import SVC

# Linear Kernel
svm_linear = SVC(kernel='linear')
svm_linear.fit(X_train, y_train)
svm_linear_accuracy = accuracy_score(y_test, svm_linear.predict(X_test))
print("SVM with Linear Kernel Accuracy: {:.4f}")

# RBF kernel
svm_rbf = SVC(kernel='rbf')
svm_rbf.fit(X_train, y_train)
svm_rbf_accuracy = accuracy_score(y_test, svm_rbf.predict(X_test))
print("SVM with RBF Kernel Accuracy: {:.4f}")

SVM with Linear Kernel Accuracy: 0.9351
SVM with RBF Kernel Accuracy: 0.9764
```

k-Nearest Neighbors (k-NN)

Uses a non-linear transformation to improve classification.

A non-parametric classifier that classifies a sample by the majority vote of its k nearest neighbors.

Best k=1, accuracy:87.2%

```
In [4]: from sklearn.neighbors import KNeighborsClassifier

best_k = 1
best_accuracy = 0

for k in range(1, 11):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    knn_accuracy = accuracy_score(y_test, knn.predict(X_test))
    if knn_accuracy > best_accuracy:
        best_k = k
        best_accuracy = knn_accuracy
    print(f"K={k}, Accuracy: {:.4f}")

print(f"Optimal k: {best_k}, Accuracy: {:.4f}")

K=1, Accuracy: 0.9720
K=2, Accuracy: 0.9642
K=3, Accuracy: 0.9713
K=4, Accuracy: 0.9699
K=5, Accuracy: 0.9781
K=6, Accuracy: 0.9698
K=7, Accuracy: 0.9687
K=8, Accuracy: 0.9678
K=9, Accuracy: 0.9674
K=10, Accuracy: 0.9658
Optimal k: 1, Accuracy: 0.9720
```

Ensemble Methods (Bagging, Random Forest, Boosting)

Ensemble learning improves accuracy by combining multiple weak classifiers.

Bagging (Bootstrap Aggregating) accuracy: 94.5%

Random Forest accuracy: 96.6%

Boosting (AdaBoost & XGBoost)

- AdaBoost accuracy: 87.2%

- **XGBoost accuracy: 97.8%**

```
In [5]: from sklearn.ensemble import BaggingClassifier, RandomForestClassifier, AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from xgboost import XGBClassifier

# Bagging
bagging = BaggingClassifier(base_estimator=DecisionTreeClassifier(), n_estimators=10)
bagging.fit(X_train, y_train)
bagging_accuracy = accuracy_score(y_test, bagging.predict(X_test))
print(f"Bagging Accuracy: {bagging_accuracy:.4f}")

# Random Forest
rf = RandomForestClassifier(n_estimators=100)
rf.fit(X_train, y_train)
rf_accuracy = accuracy_score(y_test, rf.predict(X_test))
print(f"Random Forest Accuracy: {rf_accuracy:.4f}")

# AdaBoost
ada = AdaBoostClassifier(base_estimator=DecisionTreeClassifier(), n_estimators=50)
ada.fit(X_train, y_train)
ada_accuracy = accuracy_score(y_test, ada.predict(X_test))
print(f"AdaBoost Accuracy: {ada_accuracy:.4f}")

# XGBoost
xgb = XGBClassifier(use_label_encoder=False, eval_metric='mlogloss')
xgb.fit(X_train, y_train)
xgb_accuracy = accuracy_score(y_test, xgb.predict(X_test))
print(f"XGBoost Accuracy: {xgb_accuracy:.4f}")

C:\Users\91996\AppData\Roaming\Python\Python311\site-packages\sklearn\ensemble\_base.py:166: FutureWarning: 'base_estimator' was renamed to 'estimator' in version 1.2 and will be removed in 1.4.
  warnings.warn(
Bagging Accuracy: 0.9458
Random Forest Accuracy: 0.9666
C:\Users\91996\AppData\Roaming\Python\Python311\site-packages\sklearn\ensemble\_base.py:166: FutureWarning: 'base_estimator' was renamed to 'estimator' in version 1.2 and will be removed in 1.4.
  warnings.warn(
  AdaBoost Accuracy: 0.8727
C:\Users\91996\anaconda3\lib\site-packages\xgboost\core.py:158: UserWarning: [19:04:40] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-1-08cbc033d8d4aae1-1\xgboost\xgboost-ci-windows\src\learner.cc:748:
Parameters: { "use_label_encoder" } are not used.
  warnings.warn(msg, UserWarning)
XGBoost Accuracy: 0.9781
```

Performance Comparison and Results

Accuracy and Computational Efficiency

<u>Model</u>	<u>Accuracy (%)</u>
Decision Tree	87.1%
Gaussian Naïve Bayes	55.1%
Multinomial Naïve Bayes	82.8%
SVM (Linear Kernel)	93.5%
SVM (RBF Kernel)	97.6%
k-NN (Best k=1)	97.2%

<u>Model</u>	<u>Accuracy (%)</u>
Bagging (Decision Trees)	94.5%
Random Forest	96.6%
AdaBoost	87.2%
XGBoost	97.8%

Key Observations

Best Performing Model: XGBoost (97.8%)

Most Efficient Model: Naïve Bayes

Best Trade-off Between Accuracy & Efficiency: Random Forest

Acknowledgment of Sources:

For this assignment, I consulted several resources, including Google, ChatGPT, and various online references. These resources assisted me in resolving errors, offered insights on approaching specific tasks, and greatly minimized the time required for research and troubleshooting.